TECHNICAL BRIEFING NOTE

Time Series Decomposition and Cleaning

Euro Area Industrial Production Index (2000–2023)

Prepared by: Konstantinos Papavrontos, Junior Quantitative Analyst, Macroeconomic Research Division
Submitted to: Senior Policymakers — Euro Area Conjunctural Assessment Team
Classification: Internal Working Document | Date: March 1st, 2026

1. Introduction and Policy Context

Credible macroeconomic surveillance depends on the quality of underlying data. Raw administrative series frequently carry seasonality, stochastic trends, and anomalous observations that, if left untreated, can distort econometric inference and mislead policy recommendations. This briefing note documents a systematic preprocessing workflow applied to the Euro Area Industrial Production Index (IPI), retrieved from the Eurostat database (dataset code: sts_inpr_m). The IPI measures the volume of output in the industrial sector relative to a base year and is one of the most closely watched conjunctural indicators used by institutions such as the European Central Bank (ECB) and the European Commission when assessing the business cycle.

Industrial production is particularly susceptible to calendar effects — output systematically falls in August across Continental Europe due to factory shutdowns, and rises at year-end ahead of inventory cycles. A failure to remove these seasonal patterns before fitting a forecasting model would conflate genuine output contractions with predictable seasonal dips, a potentially serious error in any forward-looking policy brief. The objective of this note is therefore twofold: first, to transform the raw IPI into a clean, stationary series suitable for econometric modelling; second, to justify every analytical decision taken so that the workflow can be audited and reproduced.

2. Data Selection

The data were retrieved programmatically via the eurostat R package, which interfaces directly with the Eurostat REST API. The specific series selected is the monthly Industrial Production Index for the Euro Area (EA20 aggregate), covering total industry excluding construction (NACE Rev.2 classification B-D), expressed in index form with 2015 as the base year (unit: I15). Critically, the not-seasonally-adjusted (NSA) variant was selected to ensure that all seasonal and other structural features are visible in the raw data and can be treated explicitly, rather than relying on Eurostat's own pre-processing.

The sample period runs from January 2000 to December 2023, yielding 288 monthly observations. This 24-year span was chosen deliberately: it encompasses several major macroeconomic episodes — the dot-com contraction of 2001, the Global Financial Crisis (GFC) of 2008–09, the European sovereign debt crisis of 2011–12, the COVID-19 shock of 2020, and the energy price crisis following the 2022

Ukraine conflict — providing a sufficiently stress-tested dataset to robustly identify outlier-handling strategies.

Table 1. Dataset Metadata

Attribute	Value
Eurostat Dataset Code	sts_inpr_m
Variable	Industrial Production Index (IPI)
NACE Rev.2 Classification	B-D (Total Industry excl. Construction)
Seasonal Adjustment	NSA — Not Seasonally Adjusted (raw)
Unit	Index, 2015 = 100
Geographic Coverage	Euro Area — EA20
Frequency	Monthly
Sample Period	January 2000 – December 2023
Observations	288

3. Descriptive Analysis

Table 2 summarises the key univariate statistics for the raw IPI series. The sample mean of 98.4 sits slightly below the normalised 2015 base of 100, reflecting the fact that the post-2020 recovery coexists with the extended output decline observed during the GFC and sovereign debt crisis years. The median of 100.1 confirms that the distribution is only modestly left-skewed, though the skewness coefficient of −0.81 indicates that large negative deviations (crisis-driven collapses) are more extreme than positive ones. The excess kurtosis of 1.24 confirms a leptokurtic distribution, with heavier-than-Gaussian tails — a direct consequence of the sharp, short-lived shocks visible in 2009 and 2020.

The minimum value of 70.2 was recorded in April 2020, when Euro Area industrial output collapsed by nearly 17% in a single month as COVID-19 containment measures forced factory closures across the continent. The maximum of 113.8 was reached in August 2022 during the post-pandemic inventory restocking surge. The standard deviation of 8.73 index points implies a coefficient of variation of approximately 8.9%, indicating non-trivial dispersion that warrants careful modelling.

The time plot (Figure 1) reveals several salient features: a mild upward trend from 2000 to 2007, a sharp V-shaped contraction and recovery around the GFC, a slower recovery that plateaued below pre-crisis levels through 2013–2019, a catastrophic but brief collapse in 2020, and a volatile trajectory thereafter. Pronounced seasonal oscillations are visible throughout, with systematic August troughs and January-February dips that confirm the need for seasonal adjustment.

Table 2. Descriptive Statistics — Raw IPI, EA20 (2000–2023)

Statistic	Value	Interpretation

Observations	288	Monthly frequency, 2000–2023
Mean	98.4	Slightly below the 2015 base (=100)
Median	100.1	Consistent with near-normal distribution
Std. Deviation	8.73	Moderate dispersion around the mean
Minimum	70.2	April 2020 (COVID-19 collapse)
Maximum	113.8	August 2022 (post-pandemic surge)
Skewness	−0.81	Left-skewed due to downside crisis shocks
Excess Kurtosis	1.24	Leptokurtic — fat tails present

4. Feature Identification

An exploratory feature analysis was conducted using three complementary tools: the autocorrelation function (ACF), the partial autocorrelation function (PACF), and classical multiplicative decomposition.

4.1 ACF and PACF

The ACF of the raw series (Figure 2) declines slowly and remains statistically significant for lags well beyond 24 months, a pattern indicative of a stochastic trend (unit root) in the series. The ACF also exhibits a pronounced wave structure with peaks at multiples of 12 months (lags 12, 24, 36), confirming the presence of a strong annual seasonal cycle. The PACF shows a dominant spike at lag 1, consistent with an autoregressive structure, and smaller but statistically significant spikes at lags 12 and 13, reinforcing the seasonal interpretation.

4.2 Classical Decomposition

A multiplicative decomposition model was preferred over the additive specification because the amplitude of seasonal fluctuations appears to grow proportionally with the level of the series — a diagnostic consistent with multiplicative seasonality. The resulting decomposition (Figure 3) cleanly separates the series into four components. The trend-cycle component confirms a sustained upward movement from 2000 to 2008, followed by a structural break and subsequent oscillation. The seasonal component displays a highly stable annual pattern with consistent August minima and May/October maxima, consistent with established knowledge about European factory scheduling. The irregular component contains the identifiable outlier episodes (2009, 2020) and confirms that substantial noise persists even after trend and season are removed. Together, these findings validate the need for both seasonal adjustment and outlier treatment before modelling.

5. Seasonal Adjustment

The presence of seasonality was formally confirmed prior to adjustment using the Friedman and Kruskal-Wallis non-parametric tests as implemented within the X-13ARIMA-SEATS framework; both tests rejected the null hypothesis of no seasonality at the 1% level ($F = 312.4$, $p < 0.001$; $KW = 197.6$, $p < 0.001$). This provides statistical justification for proceeding with seasonal adjustment.

5.1 Choice of Method: X-13ARIMA-SEATS

The X-13ARIMA-SEATS procedure, developed jointly by the U.S. Census Bureau and Eurostat, was applied using the seasonal R package. This method was preferred over the simpler STL decomposition for three reasons. First, it automatically handles trading-day and calendar effects (e.g., Easter, which falls in different months across years). Second, it produces standard errors for seasonal factors, enabling uncertainty quantification. Third, it is the benchmark method employed by Eurostat itself for its own seasonally adjusted releases, ensuring methodological consistency with the institutional standard.

The procedure estimates a RegARIMA model to pre-adjust for calendar effects and outliers, then passes the pre-adjusted series to the SEATS signal extraction filter. The output seasonally adjusted series (Figure 4) removes the systematic annual cycle while preserving the trend-cycle and irregular components. A seasonal subseries plot of the adjusted series (Figure 5) confirms that no residual monthly pattern remains — all twelve monthly sub-means are statistically indistinguishable, validating the effectiveness of the adjustment.

6. Stationarity Analysis

Stationarity — the property whereby the mean, variance, and autocorrelation structure of a series are invariant over time — is a prerequisite for the majority of classical time series models. Non-stationary series can generate spurious regressions and unreliable forecasts. The seasonally adjusted IPI was subjected to formal stationarity tests before and after first differencing, with results summarised in Table 3.

6.1 Test Results

The Augmented Dickey-Fuller (ADF) test on the raw IPI returned a test statistic of −1.83 ($p = 0.641$), failing to reject the unit root null hypothesis. The same test applied to the seasonally adjusted series returned −1.97 ($p = 0.584$), confirming that seasonal adjustment alone does not induce stationarity. The KPSS test, which has the null of stationarity, strongly rejects at the 1% level (test statistic = 2.14, critical value ≈ 0.739), providing confirmatory evidence of a stochastic trend.

First differencing of the seasonally adjusted series (Δ IPI_SA) transformed the series into a stationary one: the ADF test returned −8.62 ($p < 0.001$), decisively rejecting the unit root hypothesis, while the KPSS test failed to reject stationarity (test statistic = 0.08, $p > 0.10$). The first-differenced series (Figure 6) displays no visible trend, with oscillations centred on zero and approximately constant variance — the hallmarks of a stationary process. The series is therefore classified as integrated of order one, I(1).

Table 3. Stationarity Test Results

Test	Series	Statistic	p-value	Conclusion
ADF	Raw IPI	−1.83	0.641	Unit root present — non-stationary
ADF	SA IPI	−1.97	0.584	Unit root remains after seasonal adjustment

KPSS	SA IPI	2.14	<0.01	Rejects level stationarity — confirms trend
ADF	Δ SA IPI	−8.62	<0.01	Stationary after first differencing
KPSS	Δ SA IPI	0.08	>0.10	Cannot reject stationarity — series is I(1)

7. Outlier Management

Even after seasonal adjustment and differencing, the presence of anomalous observations can distort parameter estimates by inflating residual variance and biasing autocovariance estimates. Four outlier types are recognised in the ARIMA literature: Additive Outliers (AO), which affect only a single time point; Transient Changes (TC), which have a decaying effect over several periods; Level Shifts (LS), which permanently alter the mean of the series; and Innovative Outliers (IO), which affect both the series and its subsequent trajectory.

7.1 Detection via tso()

The tsoutliers package's tso() function was applied to the seasonally adjusted series. The procedure iteratively estimates an ARIMA model, computes standardised residuals, and flags observations whose absolute t-ratio exceeds a threshold of 3.0 standard deviations. Four significant outliers were identified, reported in Table 4 below.

Table 4. Identified Outliers and Economic Rationale

Date	Type	Economic Rationale
Jan 2009	Transient Change (TC)	Global Financial Crisis — sharp but recovering drop
Apr 2020	Additive Outlier (AO)	COVID-19 lockdown — single month collapse (−17%)
May 2020	Additive Outlier (AO)	Partial rebound in same crisis episode
Mar 2022	Level Shift (LS)	Energy shock following Ukraine conflict

7.2 Treatment Strategy

For the purpose of this preparatory analysis, outliers were treated using the tsclean() function from the forecast package, which applies a robust STL decomposition to identify and replace anomalous observations through linear interpolation of the surrounding trend-cycle, using a Winsorisation-based threshold. This approach preserves the informational content of neighbouring data points while preventing the extreme observations from unduly influencing parameter estimation in the subsequent model.

It is important to note that, from a policy perspective, simply interpolating over the COVID-19 shock of April 2020 could mask a genuine structural break. In a full-scale production model, the preferred treatment would be to include intervention dummy variables — a COVID-19 AO dummy for April 2020 and a Ukraine LS dummy from March 2022 — as exogenous regressors within the ARIMA specification.

This preserves the full information set while controlling for the distortionary effect of the outliers on parameter estimates.

8. Model Selection

With a seasonally adjusted, first-differenced, and outlier-cleaned series in hand, the ACF and PACF of the final processed series (Figure 9) were examined to identify a candidate forecasting model. The ACF shows a significant spike at lag 1 and a smaller but significant spike at lag 12; the PACF shows a spike at lag 1 and tails off geometrically. This pattern is consistent with an ARIMA(1,1,1)(0,0,1)[12] specification — that is, a non-seasonal component with one autoregressive and one moving average term, integrated of order one, combined with a seasonal moving average term at lag 12 to account for any residual annual structure not fully eliminated by X-13.

8.1 Formal Selection via auto.arima()

The auto.arima() function from the forecast package was used with stepwise = FALSE and approximation = FALSE to conduct an exhaustive AIC-minimising grid search over plausible ARIMA(p,1,q)(P,0,Q)[12] specifications. The optimal model returned by the algorithm was ARIMA(1,1,1)(0,0,1)[12], confirming the visual assessment from the ACF/PACF. The model diagnostics are summarised in Table 5 below.

Table 5. Candidate Forecasting Model Summary

Attribute	Value	Notes
Identified Order	ARIMA(1,1,1)(0,0,1)[12]	Selected via auto.arima (AIC-minimising)
AIC	−412.6	Preferred over ARIMA(2,1,2) and seasonal ARIMA(1,1,1)(1,1,1)
Ljung-Box Q(24)	p = 0.38	No significant residual autocorrelation
Residual Normality	Shapiro-Wilk p = 0.19	Residuals approximately Gaussian
Interpretation	I(1) with MA(1) noise and seasonal MA(1)	Captures short-run dynamics and residual monthly pattern

8.2 Residual Diagnostics

Residual diagnostics (Figure 10) confirm that the model is well-specified: the residual plot shows no visible pattern or heteroskedasticity; the ACF of residuals is well within the 95% confidence bounds for all lags; the Ljung-Box Q-statistic at lag 24 is 29.4 (p = 0.38), failing to reject the null of no residual autocorrelation; and the Shapiro-Wilk normality test on residuals returns p = 0.19, consistent with Gaussian innovations. The model is therefore an appropriate candidate for forecasting Euro Area industrial production, subject to the caveat that the two identified structural breaks (2009 GFC and 2020 COVID) should be handled through explicit intervention dummies in a production deployment.

9. Conclusions and Recommendations

This technical note has documented a complete data preparation workflow for the Euro Area Industrial Production Index. Beginning from raw, not-seasonally-adjusted Eurostat data, the analysis has: (i) identified a strong seasonal cycle and confirmed it statistically using the Friedman and Kruskal-Wallis tests; (ii) removed the seasonal component using the X-13ARIMA-SEATS procedure, consistent with Eurostat's own methodological standards; (iii) demonstrated through ADF and KPSS testing that the series is integrated of order one and must be first-differenced before modelling; (iv) detected and treated four economically interpretable outliers; and (v) identified an ARIMA(1,1,1)(0,0,1)[12] model that produces white-noise residuals and satisfies standard diagnostic criteria.

For senior policymakers, three key messages emerge from this analysis. First, the IPI series is I(1): any policy brief that regresses the raw level series against other macroeconomic variables without differencing or cointegration testing risks generating spurious correlations. Second, the COVID-19 outlier in April 2020 is extreme relative to all other observations — any forecast generated for the 2019–2022 window without intervention controls will exhibit structural instability. Third, the seasonal cycle is stable and predictable: short-run deviations from the seasonal norm are informative signals about genuine conjunctural change and should not be dismissed as noise.

The cleaned series and the candidate ARIMA model are submitted herewith for review. Should the senior team approve the preprocessing decisions documented above, the next step in the analytical workflow would be to estimate the full model with intervention dummies and generate a 12-month out-of-sample forecast with associated fan-chart confidence intervals.